

A Brief Introduction To R

L.Y. Deng & M.H. Garzon

Learning Objectives

- To install the package R
- To identify the basic features of R.
- To identify main data types in R, including data frames and matrices.
- To learn how to use main functions, operators, and loops.

Homepage for R

- The R package, documentation and tutorials are available for download/view at <http://www.r-project.org/> (CRAN sites)
- Manuals in pdf form are available too
- Another helpful site is at <http://statmethods.net/index.html> (Quick R)

Introduction – Why R ?

- R is a statistical programming language and environment for data manipulation, calculation and graphical display.
 - ◆ many useful operators for arrays and matrices.
 - ◆ many handy tools for interactive data analysis.
 - ◆ great graphical facilities for data analysis.
 - ◆ a programming language with conditionals, loops, user defined functions and input and output facilities.

Introduction – Features of R

- R is an interpreted computer language.
 - branching and looping as well as modular programming using functions.
 - user-defined functions in R are usually written in R, calling upon a smaller set of internal primitives.
 - allows user interface to procedures written in C, C++ or FORTRAN languages
 - for efficiency,
 - write additional primitives.

Strength of R: What R can do ?

- data handling and manipulation: numeric, textual and many matrix operations
- high-level data analytic and statistical functions
- simple to produce great graphics
- programming language: loops, branching, subroutines
- it is free and it has a strong user-support

Weakness of R

- R is not a database, but it can be connected to DBMSs
- R is basically a command-line interface but some package like `Rcmdr` can provide nice graphical user interfaces.
- R is an interpreted language which can be very slow, but you can call own C/C++ code from R.
- R lacks many spread sheet features, but R can input/output data with Excel

Data Analysis and Presentation

- The R distribution contains functionality for large number of statistical procedures.
 - ◆ linear and generalized linear models
 - ◆ nonlinear regression models
 - ◆ time series analysis
 - ◆ classical parametric and nonparametric tests
 - ◆ clustering
 - ◆ smoothing
- R also has a large set of functions which provide a flexible graphical environment for creating various kinds of data presentations.

Documentation and help file in R

- All the R functions have been documented in the form of help pages in an “output independent” form which can be used to create versions for HTML, LATEX, text etc.
 - ◆ The document “An Introduction to R” provides a more user-friendly starting point.
 - ◆ An “R Language Definition” manual
 - ◆ More specialized manuals on data import/export and extending R.

Standard packages in R

- Classical and modern statistical techniques have been implemented.
- There are several packages supplied with R (called “standard” packages) and many are available through internet sites (via <http://cran.r-project.org>).
- `install.packages()` – lists packages available to install over the internet site

Issuing commands in R

- Start R: click the icon of R after you have successfully installed the R.
- When R is started, it will prompt (`>`) and you can type in any R command.
- After you finished typing in a R command, just hit **Enter** key.
- After R finished executing your command, it will display a prompt (`>`) for your next command.
- **q()** – quits R, you will be asked whether to save workspace created.

The Workspace

- The workspace contains any user-defined objects that you might have created during an open session of R.
 - ◆ Data frames, matrices, vectors, lists
 - ◆ Functions
- Workspace is saved as a “.RData” file.
- You will want to know where your workspace is saved.

Working directory in R

- `getwd()` – displays current working directory
- `setwd("PATH")` – sets the working directory to PATH. Useful to work on different projects.
- ```
> getwd()
[1] "C:/Documents and
Settings/LYD/My Documents"
```
- ```
> setwd("C:/class/7150-2011/hw1")
```
- ```
> getwd()
[1] "C:/class/7150-2011/hw1"
```

# Storing data

---

- Every R object can be stored into and restored from a file with the commands “save” and “load”.
- This uses the XDR (external data representation) standard of Sun Microsystems and others, and is portable between MS-Windows, Unix, Mac.

```
> save(x, file="x.Rdata")
> load("x.Rdata")
```

# Managing objects in Workspace

---

- `ls()` – lists all objects currently in the workspace
- `rm()` – removes the object specified.
- ```
> ls()  
[1] "WD"
```
- ```
> rm(WD) ##or rm("WD")
```
- ```
> ls()  
character(0)
```

Command History

- You can save all the commands executed in R by saving your command history
- Click File, then click “Save History...”
- Choose directory where you want to save then click OK.
- Command history is saved in a “.RHistory” file
- `history()` lists last 25 commands
- `history(max.show=Inf)` lists all commands

Built-in dataset in R

- R has many built-in datasets that you do not have to create by yourself.
- For example, R has dataset, called **mtcars**, from 1974 *Motor Trend* US magazine, for fuel consumption (**mpg**) and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models).
- To see the list and description of the built-in datasets, type `data()`

mtcars data listing

	mpg	cyl	displacement	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

Partial listing of a dataset

- you can use `head(d,n)`, `tail(d,n)`, `print(d)` (or simply `d`) to display the first `n`, bottom `n` and all (if not too many) of the dataset `d`.

- `> head(mtcars, 2)`

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21	6	160	110	3.9	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21	6	160	110	3.9	2.875	17.02	0	1	4	4

- `> tail(mtcars, 2)`

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Maserati Bora	15.0	8	301	335	3.54	3.57	14.6	0	1	5	8
Volvo 142E	21.4	4	121	109	4.11	2.78	18.6	1	1	4	2

Special characters in R

- # #user's comment
- <- #assignment statement (also allowed: = , -> , <<- , ->>)
 - ◆ we will use **only** <- for assignment.
- [] #indexing of arrays, matrices, dataframes, lists
- () #encloses function input variables/arguments
- { } #groups statements (e.g. loops, function definition, if)
- ; #separates several statements on a single line
- \$ #extracting elements from lists or data frames
 - ◆ "\$" is similar to "." in other languages like C/C++/Java.

Variable name

- Like many modern languages (C, C++, Java), the variable names are **case-sensitive**.
- While R does **not** have a concept of “**reserved words**”, several variable/function names are better treated as of “reserved words” mainly for the purpose of readability.
 - ◆ e.g. one-letter “reserved words”: `c`, `q`, `t`, `C`, `D`, `F`, `I`, and `T`.
 - ◆ `c` (concatenate), `q`(quit) , `t`(transpose of matrix), `F`(false), `T`(true), `D`(derivative), ...

Basic data types in R

primitive (or: atomic) data types in R are:

- numeric (integer, double, complex)
- character
- logical
- function

We can build vectors, arrays, lists from basic data type.

The primary data type in R is the vector.

Operators in R

- `> x <- 2 ; y <-3`
- `> x + y`
[1] 5
- `> x * y`
[1] 6
- `> x / y` #default is floating point division
[1] 0.6666667
- `> x %/% y` # integer division
[1] 0
- `> y %/% x` # integer division
[1] 1
- `> x ^ y`
[1] 8

Useful functions on string

- `paste()`
 - ◆ # concatenate and converts to string
- `substr()`, `strsplit()`
 - ◆ # substrings and splitting strings
- `grep()`, `gsub()`
 - ◆ # finds matches, replaces matches in a string
- `tolower()`, `toupper()`
 - ◆ # uppercase, lowercase conversion
- `nchar()`
 - ◆ # number of characters in string

Example of string functions in R

- `> substr("abcdef", 2, 4)`
`[1] "bcd"`
- `> x <- "This is a"`
- `> y <- "test only"`
- `> z <- paste(x, y); z`
`[1] "This is a test only"`
- `> toupper(z)`
`[1] "THIS IS A TEST ONLY"`
- `> nchar(z)`
`[1] 19`
- `> w <- paste(z, "your score is", 90); w`
`[1] "This is a test only your score is 90"`

Concatenation and selection

- `> x <- c(2, 3, 4)`
- `> y <- c(6, 9, 2)`
- `> z <- c(x, y); z`
[1] 2 3 4 6 9 2
- `> x[c(1,3)]`
[1] 2 4
- `> x[-2]`
[1] 2 4

```
> length(z)
[1] 6
> x + y
[1] 8 12 6
> x / y
[1] 0.33333333 0.33333333
2.00000000
> x %/% y
[1] 0 0 2
```

Simple functions in R

- `> x <- c(2,3,4)`
- `> sin(x)`
[1] 0.9092974 0.1411200 -0.7568025
- `> cos(x)`
[1] -0.4161468 -0.9899925 -0.6536436
- `> sin(x)^2+cos(x)^2 #why ? all = 1`
[1] 1 1 1
- `> log(x)`
[1] 0.6931472 1.0986123 1.3862944
- `> exp(x)`
[1] 7.389056 20.085537 54.598150
- `> log10(x)`
[1] 0.3010300 0.4771213 0.6020600

Missing values and NaNs

R has some special values

- NA represents a missing value in the dataset
- NaN (**not a number**) because of the mathematical operations such as $0/0$.
- Inf (**positive infinity**) e.g. $1/0$
- -Inf (**negative infinity**) e.g. $\log(0)$
- NULL is an empty vector or array.

We can check them by

- `is.infinite(x)`
- `is.nan(x)`
- `is.na(x)`

Sequence generation in R

- Common ways to generate a sequence:
 - ◆ `from:to` # increment ± 1 .
 - ◆ `seq(from, to, by= gap)` increment or length can be specified
 - ◆ `rep(d,n)` # replicate d n times.
- ```
> x <- 9:5; x
[1] 9 8 7 6 5
```
- ```
> y <- seq(0.9,0.5, -0.1); y  
[1] 0.9 0.8 0.7 0.6 0.5
```
- ```
> z <- rep(x, 2); z
[1] 9 8 7 6 5 9 8 7 6 5
```

# Logical comparisons in R

---

- Comparing  $x$  and  $y$  (vector or scalar) with logical comparison, it will yield a vector of True/False.
  - $x < y$  ,  $x \leq y$ 
    - $\#x$  is less than, less or equal to,  $y$
  - $x > y$  ,  $x \geq y$ 
    - $\#x$  is greater than, greater or equal to,  $y$
  - $x == y$  ,  $x != y$ 
    - $\# x$  equal, not equal to,  $y$

# Logical operations in R

---

- We can use some logical operators for conditional expression:
  - `! , & , | , xor(x,y)`
    - `#` not, and, or, exclusive or
  - `any( )`
    - `#` true if any of a vector is true
  - `all( )`
    - `#` true if all values of a vector are true

# ..Logical operations in R

---

- `> x <- c(1, 5, 7, 6); y <- c(2, 6, 4, 3)`
- `> x > 3 & x < 7`  
[1] FALSE TRUE FALSE TRUE
- `> x <= y`  
[1] TRUE TRUE FALSE FALSE
- `> x[x <= y]`  
[1] 1 5
- `> (x > 3) | (y < 4)`  
[1] TRUE TRUE TRUE TRUE
- `> (x > 3)`  
[1] FALSE TRUE TRUE TRUE
- `> (y < 4)`  
[1] TRUE FALSE FALSE TRUE
- `> (x > 3) & (y < 4)`  
[1] FALSE FALSE FALSE TRUE



# Vectors and arrays

---

- **vector** is the simplest data structure used in R which is created using `c()` function.
- **array** is an ordered collection of data of the **same** type with an **integer as its index**.
  - ◆ an array can have many dimensions.
  - ◆ matrix is simply a 2-dim array.

# Using array in R

- `> x <- c(3, 5, 7, 11, 13, 19); x`  
[1] 3 5 7 11 13 19
- `> y <- array(x, dim=c(2,3)); y`  
[,1] [,2] [,3]  
[1,] 3 7 13  
[2,] 5 11 19
- `> dim(x) <- c(3,2); x`  
[,1] [,2]  
[1,] 3 11  
[2,] 5 13  
[3,] 7 19

# List in R

---

- List in R is an object consisting of a collection of objects (components) of (possibly) different types.
- The entry of the list index is usually by some names as the key.
- It can also be referenced by its position with an integer.

# Using list in R

---

- `> customer <- list(name="Fred", wife="Mary",  
+ no.children=3, child.ages=c(4,7,9))`
- `> customer$name`  
[1] "Fred"
- `> customer$child.ages`  
[1] 4 7 9
- `> customer[2]`  
\$wife  
[1] "Mary"
- `> customer[[2]]`  
[1] "Mary"

# Creating matrix

- ```
> M1 <- matrix(c(1,2,3, 11,12,13), nrow = 2, ncol=3,  
+ byrow=TRUE, dimnames = list(c("row1", "row2"),  
+ c("C.1", "C.2", "C.3")))
```
- ```
> M1
```

|      | C.1 | C.2 | C.3 |
|------|-----|-----|-----|
| row1 | 1   | 2   | 3   |
| row2 | 11  | 12  | 13  |
- ```
> M2 <- matrix(c(1,2,3, 11,12,13), nrow = 2, ncol=3,  
+ dimnames = list(c("row1", "row2"),  
+ c("C.1", "C.2", "C.3")))
```
- ```
> M2
```

|      | C.1 | C.2 | C.3 |
|------|-----|-----|-----|
| row1 | 1   | 3   | 12  |
| row2 | 2   | 11  | 13  |

# Matrix operations

- ```
> M1 + M2
```

	C.1	C.2	C.3
row1	2	5	15
row2	13	23	26
- ```
> M1 * M2 # element-wise multiplication
```

|      | C.1 | C.2 | C.3 |
|------|-----|-----|-----|
| row1 | 1   | 6   | 36  |
| row2 | 22  | 132 | 169 |
- ```
> t(M2)
```

	row1	row2
C.1	1	2
C.2	3	11
C.3	12	13
- ```
> M1 %*% t(M2) # multiplication
```

|      | row1 | row2 |
|------|------|------|
| row1 | 43   | 63   |
| row2 | 203  | 323  |

# Other matrix functions/op.

---

- `dim(A)`
  - ◆ #returns dimension of matrix or array A
- `nrow(A), ncol(A), NROW(A), NCOL(A)`
  - ◆ #number of rows and columns of matrix A
- `rownames(A), colnames(A)`
  - ◆ #names of rows and columns of matrix A
- `% * %`
  - ◆ # matrix multiplication

# Other matrix functions/op.

---

- `t(A)` # transpose of matrix A
- `solve(A)` # inverse of matrix A
- `svd(A)` , `qr(A)` , `chol(A)`
  - ◆ # singular value, QR, cholesky decomposition of matrix A
- `eigen(A)` , `det(A)`
  - ◆ # eigenvalues and eigenvectors, determinant of matrix A



# Combining matrices and arrays

- `cbind(x, y)`
- # binds matrices, dataframes, ... columnwise
- `rbind(x, y)`
- # binds matrices, dataframes, ... rowwise

```
> x <- c(1, 2, 7, 9); y <- 5:8
> cbind(x, y)
 x y
[1,] 1 5
[2,] 2 6
[3,] 7 7
[4,] 9 8
> rbind(x, y)
 [,1] [,2] [,3] [,4]
x 1 2 7 9
y 5 6 7 8
> c(x, y)
[1] 1 2 7 9 5 6 7 8
```

# Data frames

---

**data frame** is a rectangular table with rows and columns; data within each column has the same type (e.g. number, text, logical), but different columns may have different types.

**data.frame():**

- an R command to create data frames, tightly coupled collections of variables which share many of the properties of matrices and of lists,
- used as the fundamental data structure by most of R's modeling software.

# Creating data frame

---

- You can recreate a data frame from scratch by  
`my_data <- edit(data.frame())`  
that you can enter data into the given form.
- You can also import from external file (to be discussed later) or you can save the data created.
- `> my_data <- data.frame(x,y); my_data`

|   | x | y |
|---|---|---|
| 1 | 1 | 5 |
| 2 | 2 | 6 |
| 3 | 7 | 7 |
| 4 | 9 | 8 |

# Data subsetting in R

---

- `x[n], x[-n]` # select nth element, all but nth element from vector x
- `x[1:n], x[-(1:n)]` # select first n elements, all but first n elements from x
- `x[c(1,4,6)]` # select element 1,4 and 6 from vector x
- `x[x>3 & x<5]` # select elements that meet condition
- `which(x==3)` # returns indices to values x that meet the condition

# ..Data subsetting in R

---

- `> x <- c(2, 5, 7, 11, 13, 17)`
- `> x[3]`  
`[1] 7`
- `> x[-3]`  
`[1] 2 5 11 13 17`
- `> x[1:3]`  
`[1] 2 5 7`
- `> x[-(1:3)]`  
`[1] 11 13 17`
- `> x[c(1,4,6)]`  
`[1] 2 11 17`
- `> x[-c(1,4,6)]`  
`[1] 5 7 13`
- `> which(x==13)`  
`[1] 5`

# Subsetting matrix/data frame in R

---

- Same rule for vector subsetting can be used for matrix or data frame (to be discussed later)
- `A[i, j], A[, j], A[i, ]` # selects element i,j, the jth column, i-th row from matrix A
- `A[, cols]` # selects columns **cols** from matrix A
- `A["name", ]` # selects row named **"name"** from matrix A
- `D$name, D[["name"]]` # selects column named **"name"** from data frame D

# ..Subsetting matrix/data frame in R

---

- `> mtcars[1:4,1:5]`

|        |         | mpg  | cyl | disp | hp  | drat |
|--------|---------|------|-----|------|-----|------|
| Mazda  | RX4     | 21.0 | 6   | 160  | 110 | 3.90 |
| Mazda  | RX4 Wag | 21.0 | 6   | 160  | 110 | 3.90 |
| Datsun | 710     | 22.8 | 4   | 108  | 93  | 3.85 |
| Hornet | 4 Drive | 21.4 | 6   | 258  | 110 | 3.08 |
- `> mtcars[1:4,1]`  
`[1] 21.0 21.0 22.8 21.4`
- `> mtcars[1:4, "mpg"]`  
`[1] 21.0 21.0 22.8 21.4`
- `> mtcars$mpg[1:4]`  
`[1] 21.0 21.0 22.8 21.4`

# if statement in R

---

- `if, else, else if`
  - ◆ `#conditionally execute statements`
  - ◆ `#useful only when comparing two values, not two vectors (why not?)`
  - ◆ `# often used with all() or any()`
- `R example:`
- `if(all(x < 0)) cat("all x values are negative\n")`



# ifelse statement in R

---

- `ifelse(cond, yes, no)`
  - ◆ `# if (component-wise) condition is true/false, executes (component-wise) statement 'yes'/'no'`
- R example:
  - `x <- c(6:-4)`
  - `sqrt(x)` #- gives warning
  - `sqrt(ifelse(x >= 0, x, NA))` # no warning

# Repetitive execution

---

- `for (el in seq) {expr}`
  - ◆ `#repeat expr for each element in seq`
- `while (cond) {expr}`
  - ◆ `#repeat expression while condition is true`
  - ◆ `#be very careful for vector comparison`
- `repeat {expr}`
  - ◆ `#repeat until break encountered`

# Breaking repetitive execution

---

- break
  - ◆ it terminates execution of for, while, repeat loops
  - ◆ it can be used to terminate any loop, possibly abnormally.
- next
  - ◆ it transfers execution to next iteration in loops
  - ◆ it can be used to discontinue one particular cycle and skip to the “next”.

# Example

```
> s <- 0
> for(i in 1:4) {
 s <- s+ i^0.5
 print(s)
}
[1] 1
[1] 2.414214
[1] 4.146264
[1] 6.146264
```

```
> i <- 1; s <- 0;
> while(s<=10) {
 s <- s+ i^0.5
 print(s); i <- i+1
}
[1] 1
[1] 2.414214
[1] 4.146264
[1] 6.146264
[1] 8.382332
[1] 10.83182
```

# User-defined functions

---

## Example:

```
f <- function(a, b)
{
 return (a+b)
}
```

## Note:

- Note that `return` is a function in R; its argument must be contained in parentheses.
- The use of `return` is optional; otherwise the value of the last line executed in a function is its return value.

# apply( arr, margin, fct )

---

Apply the function fct along some dimensions of the array arr, according to margin, and return a vector or array of the appropriate size.

```
> x
 [,1] [,2] [,3]
[1,] 5 7 0
[2,] 7 9 8
[3,] 4 6 7
[4,] 6 3 5
> apply(x, 1, sum)
[1] 12 24 17 14
> apply(x, 2, sum)
[1] 22 25 20
```

# lapply(li, *function*)

---

- To each element of the list *li*, the function *function* is applied.

```
> li <- list("This", "example", "is", "great")
```

```
> lapply(li, toupper)
```

```
[[1]]
```

```
[1] "THIS"
```

```
[[2]]
```

```
[1] "EXAMPLE"
```

```
[[3]]
```

```
[1] "IS"
```

```
[[4]]
```

```
[1] "GREAT"
```

# apply( li, fct )

apply is a simplified version of lapply by default returning a vector or matrix if appropriate

```
> li <- list("This", "example", "is", "great")
> apply(li, toupper)
 [1] "THIS" "EXAMPLE" "IS" "GREAT"
> fct = function(x) { return(c(x, x*x, x*x*x)) }
> apply(1:5, fct)
 [,1] [,2] [,3] [,4] [,5]
[1,] 1 2 3 4 5
[2,] 1 4 9 16 25
[3,] 1 8 27 64 125
```



# Input output in R

---

- By default, the input is from the keyboard and output is to the screen. However, there are many other methods can be used.
- `write.table(x, file)`
  - ◆ # writes object x as a dataframe to a table
- `read.table(file)`
  - ◆ # reads table from space-delimited file, aligned in columns
- `read.csv(file), read.delim(file)`
  - ◆ # reads table comma- delimited or tab-delimited file

# Reading data from files

---

| Price | Floor | Area | Rooms | Age | Cent.heat |
|-------|-------|------|-------|-----|-----------|
| 52.00 | 111.0 | 830  | 5     | 6.2 | no        |
| 54.75 | 128.0 | 710  | 5     | 7.5 | no        |
| 57.50 | 101.0 | 1000 | 5     | 4.2 | no        |

...

```
HousePrice <-
read.table("houses.data", header=TRUE)
```

# Importing and exporting data

---

There are many ways to get data into R and out of R.

Most programs (e.g. Excel), as well as humans, know how to deal with rectangular tables in the form of tab-delimited text files.

```
> x = read.delim("filename.txt")
```

also: `read.table`, `read.csv`

```
> write.table(x, file="x.txt", sep="\t")
```

# Importing data

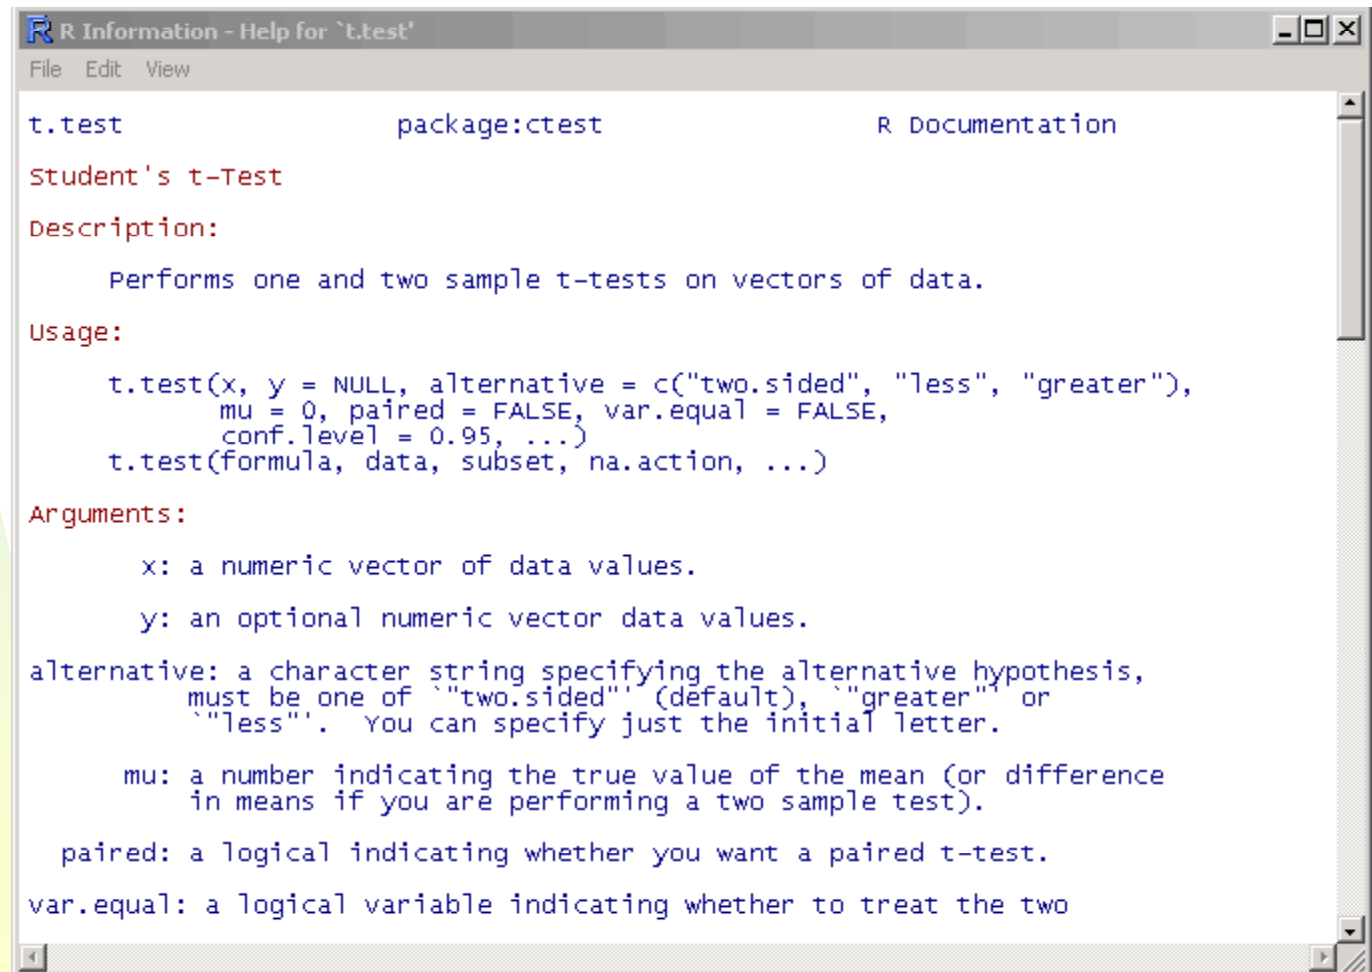
---

- Type conversions: by default, the read functions try to guess and autoconvert the data types of the different columns (e.g. number, factor, character).
  - There are options `as.is` and `colClasses` to control this – *read the online help*
  - Understand the conventions your input files use and set the quote options accordingly.

# Getting help

Details about a specific command whose name you know (input arguments, options, algorithm, results):

```
>? t.test
or
>help(t.test)
```

A screenshot of the R help window titled "R Information - Help for 't.test'". The window has a menu bar with "File", "Edit", and "View". The main content area displays the following information:

t.test package:cctest R Documentation

Student's t-Test

Description:

Performs one and two sample t-tests on vectors of data.

Usage:

```
t.test(x, y = NULL, alternative = c("two.sided", "less", "greater"),
 mu = 0, paired = FALSE, var.equal = FALSE,
 conf.level = 0.95, ...)
t.test(formula, data, subset, na.action, ...)
```

Arguments:

x: a numeric vector of data values.

y: an optional numeric vector data values.

alternative: a character string specifying the alternative hypothesis, must be one of "two.sided" (default), "greater" or "less". You can specify just the initial letter.

mu: a number indicating the true value of the mean (or difference in means if you are performing a two sample test).

paired: a logical indicating whether you want a paired t-test.

var.equal: a logical variable indicating whether to treat the two

# Further Topics

---

Some of the topics listed will be discussed in the later modules.

- Graphics in R (page 25, much more on Chapter 15)
- Lattice graphics ( page 30, skip)
- Finer graphic parameter settings (page 27)
- Customized options setting (page 34)